

Multi-Source RAG for Technical Support

TaskFlow Lite (Fictional Product)

Srutileka Suresh

February 2026

1 Problem Statement

Build a Retrieval-Augmented QA system that answers customer support questions about a fictional software product (TaskFlow Lite) by retrieving evidence from **three distinct sources**: product documentation, customer forums, and technical blog posts. The system must implement source-specific chunking, multi-source retrieval with weighting, reranking, contradiction handling, and logging.

2 Product Overview: TaskFlow Lite

TaskFlow Lite is a lightweight project/task management tool. Users can manage projects (with a task limit), attach files to tasks (with file type + size limits), and use an API (with key expiry and rate limits). The knowledge sources intentionally include cases where community forums or blogs conflict with official docs.

3 Dataset (Custom made)

- **Documentation (.md):** Authoritative product rules such as attachment limits, export formats, API key expiry, and error codes.
- **Customer Forums(.jsonl):** User questions and replies; may be noisy or contradictory.
- **Technical Blogs(.md):** How-to posts and best practices; sometimes outdated.

4 Chunking Strategy

Different sources have different structure, so chunking is tuned per source:

- **Docs:** Larger, structured chunks to preserve bullet lists (limits, error codes) as a single coherent unit.
- **Blogs:** Medium chunks to keep narrative explanations together.
- **Forums:** Small chunks per question or per reply to keep distinct opinions separated (and to support contradiction detection).

5 Retrieval (Multi-Source + Weighting)

5.1 Why separate retrieval per source?

Three separate FAISS indexes are built (`docs`, `blogs`, `forums`). At query time, the system retrieves top- k from each index to ensure every source can contribute evidence.

5.2 Source weighting

After retrieving from each source, the system applies source weights (docs prioritized):

$$\text{weighted_score} = \text{similarity_score} \times w_{\text{source}}$$

Weights used in the run: `docs=1.2, blogs=1.0, forums=0.8` :contentReference[oaicite:2]index=2.

6 Reranking (Cross-Encoder)

The system uses a two-stage ranking pipeline:

1. **Bi-encoder retrieval** (fast): retrieve a shortlist of candidates via embedding similarity (FAISS).
2. **Cross-encoder reranking** (precise): rerank only the top candidates

This design balances speed (vector search) and relevance (token-level interaction in cross-encoder).

7 Contradiction Handling

Contradictions can occur between sources (e.g., size limits, expiry durations, rate limits). The system:

- extracts numeric claims (e.g., 25MB, 60 days, 300/hour) from top evidence,
- flags conflicts when multiple values appear for the same category,
- resolves by a reliability policy: **Docs > Blogs > Forums**.

8 Logging

Each query is logged to a JSONL file including:

- query string,
- top candidates pre- and post-rerank,
- sources used and scores,
- detected contradictions and the selected reliability policy.

9 Performance Analysis

We compare retrieval quality before reranking (`pre`) vs after reranking (`post`) using file-level “gold” targets for the synthetic dataset :contentReference[oaicite:4]index=4.

9.1 Summary Metrics

From `metrics.json` :contentReference[oaicite:5]index=5:

Metric	Pre-rerank	Post-rerank
Hit@5	0.80	1.00
Hit@1	0.70	0.40

Table 1: Retrieval performance before vs after reranking (10 queries).

Interpretation:

- Hit@5 improves from 0.80 to 1.00, meaning reranking + multi-source merging increased the chance that the correct file appears in the top-5.
- Hit@1 decreases from 0.70 to 0.40 in this run. This happens because the cross-encoder sometimes ranked short, direct forum/blog snippets above the authoritative doc chunk (even though the doc chunk remained in the top results). The system mitigates this via a **source reliability policy** during answer selection.

10 Query Result Screenshots



Figure 1: Placeholder grid for 10 screenshot outputs (replace images in `figs/`).

11 Limitations and Future Work

- The contradiction detector focuses on numeric claims (MB/day/hour). More general contradiction detection would require deeper semantic modeling.
- Hit@1 decreased after reranking in this run :contentReference[oaicite:17]index=17; a refinement is to incorporate source priors directly into reranking or to select the highest reranked chunk within the preferred source.

- The answer synthesis is deterministic (no external LLM). Adding a small LLM could improve readability, but would introduce nondeterminism and potential hallucinations.